

AYYANAR JEYAKRISHNAN

Engineering Leader | Technology Delivery, Platform Engineering & AI Architecture

jayyanar@gmail.com | +91-9600068932 | linkedin.com/in/jayyanar | dataopslabs.com | Bengaluru, India | Wharton Executive Education
AWS AI Hero | AWS 15x Certified | AWS Subject Matter Expert | TOGAF Level 2 | Azure & Google Cloud Certified | Snowflake Squad

EXECUTIVE SUMMARY

Engineering leader with 20+ years in software engineering, technology delivery and application development, including 10+ years leading technical teams and 10+ years of direct people leadership across banking, financial services, healthcare technology and enterprise cloud. Currently Executive Director – AI Solutions Engineering at Wells Fargo, leading a cross-functional engineering organization of software engineers, quality engineers, platform engineers and Scrum teams delivering the Enterprise Data Intelligence Platform (EDIP) — the enterprise platform through which multiple lines of business consume Generative AI in production.

Proven at building high-performing engineering organizations, assessing capability gaps and running structured training, mentoring and development programs; and at driving execution of technology roadmaps through Agile, DevOps, CI/CD and quality engineering with predictable, measurable delivery. Extensive experience delivering solutions built on APIs, enterprise integrations, event-driven and distributed systems on AWS, Azure and GCP, and coordinating delivery across multiple technology domains, external vendors, cloud partners and third-party engineering organizations. Defines AI/ML and GenAI architectural patterns, agentic architectures and intelligent automation capabilities, and champions enterprise-wide adoption of AI-assisted development. Influences outcomes across teams and partner organizations without direct authority; TOGAF Level 2 certified and AWS AI Hero.

CORE COMPETENCIES

People Leadership & Talent Development: Team Building & Coaching | Skills Assessment & Capability Gap Analysis | Mentoring & Development Plans | Workforce & Capacity Planning | Hiring, Onboarding & Succession | Performance Management | Engineering Culture & Inclusion

Delivery Leadership & Execution: Technology Roadmap Execution | Backlog Prioritisation & Release Planning | Sequencing & Dependency Management | Risk, Issue & Roadblock Mitigation | Predictable, Measurable Delivery | Budget & Vendor Management | Executive Stakeholder Communication

Cross-Domain & Partner Collaboration: Multi-Domain Delivery Coordination | Cross-Functional Dependency Management | External Partner, Vendor & Third-Party Engineering Collaboration | Offshore/Onshore Distributed Delivery | Influence Without Authority | Senior Governance Forums

Technical Leadership & Architecture: APIs & Enterprise Integration | Event-Driven Architecture | Distributed Systems at Scale | Cloud-Native & Serverless (AWS, Azure, GCP) | Enterprise Data Platforms | TOGAF Enterprise Architecture | Scalable, Secure & Reliable Solution Design

AI, GenAI & Intelligent Automation: AI/ML Architectural Patterns & Standards | Generative AI Platform Engineering | LLM Integration | Agentic Architectures & Multi-Agent Orchestration | Intelligent Automation | AI-Assisted Development Adoption | Spec-Driven Development

Agile, DevOps & Quality Engineering: Agile at Scale & Scrum | CI/CD & DevSecOps | Test Automation & Agentic QA | Infrastructure as Code | Observability & Site Reliability | Engineering Productivity Tooling | Security-First Design

SELECTED LEADERSHIP ACHIEVEMENTS

- Delivered \$5M+ in annual savings through an Agentic QA Platform built on multi-agent AI workflows — recognised as a Top 10 CxO-level critical programme.
- Scaled a Presto-based distributed analytics platform from zero to 5,000 concurrent users on 800+ node AWS infrastructure, enabling self-service access across business teams.
- Led data centre to AWS migration for a healthcare technology group, refactoring 60% of applications to a serverless-first architecture in partnership with external AWS delivery partners.
- Engineered a Tier-1, globally distributed, low-latency in-memory post-trade platform (4TB) sustaining 99.9%+ availability for global trading operations.
- Built and led engineering teams of 20+ across multiple geographies; established a Train-the-Trainers programme on Agentic AI upskilling engineering teams enterprise-wide.

PROFESSIONAL EXPERIENCE

Executive Director – AI Solutions Engineering (Engineering Leader)

Wells Fargo, COO Technology | Bengaluru, India | Feb 2023 – Present

Application Owner, Enterprise Data Intelligence Platform (EDIP) | LoB-wide AI-Driven Development Lifecycle Champion

- Lead a cross-functional engineering organisation spanning multiple Scrum teams of software engineers, quality engineers, platform engineers and collaborate with technical product partners, delivering and supporting the Enterprise Data Intelligence Platform (EDIP) consumed by business units across multiple lines of business.
- Own technical delivery, resourcing strategy, capacity planning, budget and vendor management for the AI engineering organisation, consistently meeting financial targets and efficiency goals.
- Assess team capabilities, identify skill gaps and run structured training, mentoring and development plans; established an enterprise Train-the-Trainers programme on Agentic AI to build organisational readiness for future technology needs.
- Mentor Principal Engineers and foster a culture of accountability, continuous learning and engineering excellence, championing next-generation practices including spec-driven development, AI-augmented workflows and systematic refactoring of brownfield estates.
- Define AI/ML and Generative AI architectural patterns and enterprise standards for how lines of business integrate LLMs into production applications, covering agentic architectures, multi-agent orchestration and intelligent automation.
- Drive enterprise-wide adoption of AI-assisted development tooling across engineering teams — code generation and review, automated test creation and maintenance, documentation and debugging — with clear expectations that engineers validate and own AI-generated output.
- Built and deployed the Agentic QA Platform, automating quality engineering through multi-agent AI workflows and delivering \$5M+ annual savings; recognised as a Top 10 CxO-level critical programme.
- Delivered DART (Document Analysis Response Tool), an enterprise AI platform providing inline translation with layout preservation, intelligent document extraction and automated classification at scale.
- Partner with product, architecture and business stakeholders to translate business objectives into executable delivery plans; drive backlog prioritisation, release planning and sequencing while balancing business value, technical complexity, dependencies and capacity.
- Identify, communicate and mitigate delivery risks, dependencies and roadblocks; report progress, milestones, implementation status and dependency impacts to senior leadership and governance forums.
- Authored 3 enterprise whitepapers selected in the Top 20 org-wide (Adaptive Synthetic Data for Fraud Detection; Text2SQL Agent at Enterprise Scale; AML Detection with Multi-Agent Orchestration).

DevOps Consultant – Capability Lead (Engineering Leader)

PSL Group (Healthcare Technology) | Canada | Aug 2021 – Feb 2023

- Built and led a team of 20+ DevOps and platform engineers, establishing the reference architecture and engineering standards for DevOps and DevSecOps on AWS, with observability (New Relic), Infrastructure as Code (CloudFormation, Terraform) and security-first design.
- Defined and executed the data centre to AWS migration strategy, refactoring 60% of applications to a serverless-first architecture, coordinating delivery across internal engineering teams and external AWS partner organisations.
- Established organisation-wide cloud adoption standards and policies jointly with Security Architects, ensuring solutions aligned to enterprise architecture, integration patterns and technology roadmaps.
- Built a serverless, event-driven data processing platform for clickstream analytics using AWS Managed Kafka (MSK), Kinesis and Kinesis Firehose, delivering real-time business insight.
- Managed cross-domain dependencies, sequencing and implementation risk across application, infrastructure, security and partner teams to deliver end-to-end migration outcomes.

Lead Data Platform Engineer

Grab Financial Services | Bengaluru, India | Apr 2019 – Jul 2021

- Built and operated the GrabPay data platform — high-volume streaming ingestion and processing on EMR Spark and Kafka with a Presto distributed query engine serving 5,000+ users across 800+ node AWS infrastructure.
- Scaled the Presto analytics platform from zero to 5,000 concurrent users, enabling self-service data access for business intelligence teams across the organisation.
- Managed the end-to-end ML platform — JupyterLab for data scientists, batch orchestration via Airflow, and model inference services exposed as APIs on EKS.
- Hired, trained and developed a high-performing data engineering team, fostering technical excellence and continuous improvement.

Data Science Systems Engineer

PSL Group (Healthcare Technology) | Canada | Sep 2018 – Apr 2019

- Engineered Snowflake Streaming Snowpipe pipelines for real-time survey analytics, transforming millions of records into actionable business intelligence.
- Partnered with data science teams to architect and deploy 20+ ML models on AWS with continuous evaluation frameworks.

Assistant Vice President – Big Data Engineering

JP Morgan Chase | India | Jun 2017 – Aug 2018

- Delivered enterprise Big Data streaming applications using Kafka, Spark Streaming and Elastic Stack on Cloudera, enabling real-time processing for critical banking operations.
- Led migration of Oracle and Cloudera Hadoop workloads to a Snowflake data warehouse on S3, optimising cost-per-query for enterprise analytics.
- Contributed to the Big Data & ML Pattern Team and led training on productionalising ML models and cloud technologies across engineering teams.

Technical Manager – Big Data

Société Générale | India | Sep 2014 – Feb 2017

- Led solution design and delivery of an enterprise Data Lake on HortonWorks, managing a team of 20+ engineers across multiple workstreams and technology domains.
- Engineered a Tier-1 post-trade application — high-throughput, low-latency, globally distributed in-memory system (Gemfire) — with CI/CD via Jenkins, Ansible and Puppet.
- Owned the Tier-1 integration and middleware estate (WebSphere, WebLogic, IBM MQ, RabbitMQ, F5, Elasticsearch, Arcot SSO), sustaining 99.9%+ uptime for global trading operations.

Advisory IT Specialist – Cloud

IBM | India | Sep 2009 – Sep 2014

- Managed IBM SoftLayer and AWS public cloud infrastructure for large-scale e-commerce applications on WebSphere Commerce, delivering end-to-end solutions for enterprise clients.
- Led a team of 14 engineers through offshore transition for Bank of Ireland and WellPoint (US), managing knowledge transfer, external stakeholder alignment and service continuity.
- Implemented IBM Hadoop on Pure Application System, led cross-data-centre high-availability deployments for clients including Bombardier, and received IBM internal patent recognition for WebSphere automation tooling.

Earlier Career – Technical Lead / Systems Engineer

Satyam (Merrill Lynch) | HCL Infosystems | Naturosoft | India & USA | Apr 2005 – Apr 2009

- Technical Lead at Satyam Computers for Merrill Lynch (USA) owning application server infrastructure and enterprise middleware operations; earlier managed 200+ node infrastructure across Active Directory, Red Hat Linux, Apache and Tomcat.

TECHNICAL SKILLS

AI & GenAI: Generative AI application platforms, LLM integration, agentic and multi-agent architectures, intelligent document processing, Text2SQL agents, synthetic data generation, MLOps/LLMOps, AI-assisted development tooling

Cloud & Platform: AWS (15x certified, AI Hero, SME), Azure, Google Cloud, EKS/Kubernetes, serverless, Lambda, EMR, MSK, Kinesis, Snowflake

Integration & Distributed Systems: REST APIs, event-driven architecture, Kafka, Kinesis, IBM MQ, RabbitMQ, WebSphere, WebLogic, F5, Gemfire in-memory grid, Spark, Presto, Airflow

DevOps & Quality: CI/CD, DevSecOps, Jenkins, Terraform, CloudFormation, Ansible, Puppet, automated and agentic test engineering, New Relic, Elastic Stack

Architecture & Practice: TOGAF Level 2, enterprise architecture standards, reference architectures, Agile/Scrum at scale, engineering governance

THOUGHT LEADERSHIP & COMMUNITY IMPACT

Blog: dataopslabs.com | GitHub: github.com/jayyanar

- Speaking: 50+ technical talks on AWS, Azure, MLOps, Generative AI and LLMOps at international conferences including AWS re:Invent, AWS User Groups, Microsoft Reactor and AZConf.
- Workshops & Publications: Led a 4-hour hands-on GenAI workshop (AWS) for 150+ participants and delivered Azure MLOps for DevOps Engineers to 125+ attendees; 3 enterprise GenAI whitepapers (Top 20 enterprise-wide), 1 patent accepted, 1 IEEE paper.